

FINISHED

WAT files - understanding the JSON structure

Took 0 sec. Last updated by anonymous at August 28 2017, 1:52:14 PM.

FINISHED

```
%pyspark

import boto
from boto.s3.key import Key
from gzipstream import GzipStreamFile
from pyspark.sql.types import *
import warc

import json

watlist = sc.textFile("s3://commoncrawl/crawl-data/CC-MAIN-2017-04/wat.paths.gz")
watlist.cache()

conn = boto.connect_s3(anon=True, host='s3.amazonaws.com')
bucket = conn.get_bucket('commoncrawl')

def unpack(uri):
    key_ = Key(bucket, uri)
    file_ = warc.WARCFile(fileobj=GzipStreamFile(key_))
    return file_

def mapper(id_, iterator):
    for uri in iterator:
        file = unpack(uri)
        for record in file:
            try:
                yield record['Content-Type']
            except KeyError:
                yield None

nfiles = 16
files = sc.parallelize(watlist.take(nfiles))

ct = files.mapPartitionsWithIndex(mapper)
ct.cache()
print(ct.count())
print(ct.countByValue())
ct.unpersist()

2621630
defaultdict(<type 'int'>, {'application/warc-fields': 16, 'application/json': 2621614})
PythonRDD[131] at RDD at PythonRDD.scala:48
```

Took 1 min 18 sec. Last updated by anonymous at August 28 2017, 12:28:00 PM.

Let's examine a sample of json records:

READY

```
%pyspark
```

FINISHED

```
from pprint import pprint

def json_mapper(id_, iterator):
    conn = boto.connect_s3(anon=True, host='s3.amazonaws.com')
    bucket = conn.get_bucket('commoncrawl')

    for uri in iterator:
        key_ = Key(bucket, uri)
        file_ = warc.WARCFile(fileobj=GzipStreamFile(key_))

        for record in file_:
            if record['Content-Type'] == 'application/json':
                record = json.loads(record.payload.read())

                try:
                    yield record
                except KeyError:
                    yield None

nrecords = 100
sample = files.\
    mapPartitionsWithSplit(json_mapper).\
    take(nrecords)

pprint(sample[1])
```

```
4ZH356NNFQZMSVX/XZKYAYSCLX5QBYJ',
u'Entity-Length': u'0',
u'Entity-Trailing-Slop-Byte',
s': u'0',
u'Headers': {u'Accept': u't
ext/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8',
u'Accept-Encoding',
ing': u'x-gzip, gzip, deflate',
u'Host': u'03o
nline.com',
u'User-Agent':
u'CCBot/2.0 (http://commoncrawl.org/faq/)'},
u'Headers-Length': u'211',
u'Request-Message': {u'Meth
od': u'GET',
u'Pat
h': u'/news/3383',
u'Venc
```

Took 1 sec. Last updated by anonymous at August 28 2017, 12:29:33 PM.

We see that the field ['Envelope']['WARC-Header-Metadata']['WARC-Target-URI'] contains the URI of the current web page:

Took 0 sec. Last updated by anonymous at August 28 2017, 12:31:20 PM.

```
%pyspark
```

FINISHED

```
for rec in sample:
    try:
        print(rec['Envelope']['WARC-Header-Metadata']['WARC-Target-URI'])
```

```
except KeyError:
    print("")
```

```
http://103.2.132.3/newsdetail.asp?ID=44
http://103.2.132.3/newsdetail.asp?ID=44
http://103.2.132.3/newsdetail.asp?ID=44
http://1037theloon.com/tags/w-a-s-p/
http://1037theloon.com/tags/w-a-s-p/
http://1037theloon.com/tags/w-a-s-p/
http://1061thetwister.iheart.com/articles/entertainment-news-104651/trick-pony-reunites-following-sevenyear-split-12059522/
http://1061thetwister.iheart.com/articles/entertainment-news-104651/trick-pony-reunites-following-sevenyear-split-12059522/
http://1061thetwister.iheart.com/articles/entertainment-news-104651/trick-pony-reunites-following-sevenyear-split-12059522/
http://1079thebear.iheart.com/onair/ken-dashow-32036/epic-fails-for-july-2014-12639622/
http://1079thebear.iheart.com/onair/ken-dashow-32036/epic-fails-for-july-2014-12639622/
http://1079thebear.iheart.com/onair/ken-dashow-32036/epic-fails-for-july-2014-12639622/
http://107jamz.com/events-lake-charles/icd-10-cm-coding-bootcamp/24-july-2013/
http://107jamz.com/events-lake-charles/icd-10-cm-coding-bootcamp/24-july-2013/
http://107jamz.com/events-lake-charles/icd-10-cm-coding-bootcamp/24-july-2013/
```

Took 0 sec. Last updated by anonymous at August 28 2017, 12:31:23 PM.

Later we'll want to aggregate records by web domain, and use the information in the individual **FINISHED** records to build features of the domains.

Let's build a traverse function to output a lists of keys of a json record together with its tree depth and boolean is-leaf indicator:

Took 0 sec. Last updated by anonymous at August 28 2017, 12:33:56 PM.

```
%pyspark FINISHED

from __future__ import print_function

def traverse(js, depth, keys):
    if type(js) is dict:
        d = depth + 1
        for k in js.keys():
            if type(js[k]) is not dict: leaf = 1
            else: leaf = 0
            keys += [(d,leaf,k)]
            depth, keys = traverse(js[k], d, keys)
    return depth, keys

js = sample[1]
depth, keys = traverse(js, 0, [])
keys.sort()

print(len(keys), "json keys:")
for x in keys:
    print(x)
(3, 1, u'WARC-Date')
(3, 1, u'WARC-IP-Address')
(3, 1, u'WARC-Record-ID')
(3, 1, u'WARC-Target-URI')
(3, 1, u'WARC-Type')
(3, 1, u'WARC-Warcinfo-ID')
(4, 0, u'Headers')
(4, 0, u'Request-Message')
(4, 1, u'Entity-Dictionary')
```

```
(4, 1, u'Entity-Digest')
(4, 1, u'Entity-Length')
(4, 1, u'Entity-Trailing-Slop-Bytes')
(4, 1, u'Headers-Length')
(5, 1, u'Accept')
(5, 1, u'Accept-Encoding')
(5, 1, u'Host')
(5, 1, u'Method')
(5, 1, u'Path')
(5, 1, u'User-Agent')
```

Took 0 sec. Last updated by anonymous at August 28 2017, 12:33:59 PM.

This json record has 41 keys. Let's see how that varies over all the records:

READY

%pyspark

FINISHED

```
from collections import Counter
import matplotlib.pyplot as plt


def get_json_keys(id_, iterator):
    for uri in iterator:
        file = unpack(uri)
        for record in file:
            if record['Content-Type'] == 'application/json':
                record = json.loads(record.payload.read())

                try:
                    _,k = traverse(record, 0, [])
                    """
                    yield the shape of the record:
                    """
                    yield dict(Counter([x[0:2] for x in k]))
                except KeyError:
                    yield None

json_keys = files.mapPartitionsWithIndex(get_json_keys)
json_shape = json_keys.collect()

total_shape = [sum(x.values()) for x in json_shape]

plt.hist(total_shape, bins=30)
plt.yscale('log')
plt.show()
```



Took 5 min 54 sec. Last updated by anonymous at August 28 2017, 12:40:16 PM.

Let's break down this histogram by tree depth and leaf vs node:

READY

```
%pyspark
from __future__ import print_function

maxdepth = max([k[0] for y in json_shape for k in y.keys()])

for depth in range(1, 1+maxdepth):
    print("Depth", depth)

    nodeshape = dict(Counter([x[(d,l)] for x in json_shape for (d,l) in x.keys() if d==depth])
    if len(nodeshape.items()) > 0:
        print("nodes:")
        for i in nodeshape.items(): print(i)

    leafshape = dict(Counter([x[(d,l)] for x in json_shape for (d,l) in x.keys() if d==depth and l==0])
    if len(leafshape.items()) > 0:
        print("leaves:")
        for i in leafshape.items(): print(i)
```

```
(38, 158)
(39, 1595)
(40, 129)
(41, 297)
(42, 573)
(43, 164)
(44, 13)
(45, 68)
(46, 8)
(54, 18)
(56, 1)
```

Depth 6

FINISHED

leaves:

```
(1, 6645)
(2, 10224)
(3, 33612)
(4, 727800)
```

Took 1 min 8 sec. Last updated by anonymous at August 28 2017, 12:41:52 PM.

In other words:

FINISHED

At depth 1:

all records have two nodes

```
(1, 0, u'Container')
(1, 0, u'Envelope')
```

At depth 2:

all records have 3 nodes and 7 leaves

```
(2, 0, u'Gzip-Metadata')
(2, 0, u'Payload-Metadata')
(2, 0, u'WARC-Header-Metadata')
(2, 1, u'Actual-Content-Length')
(2, 1, u'Block-Digest')
(2, 1, u'Compressed')
(2, 1, u'Filename')
(2, 1, u'Format')
(2, 1, u'Offset')
(2, 1, u'WARC-Header-Length')
```

At depth 3:

all records a single 3 node

```
(3, 0, u'HTTP-Request-Metadata')
```

but 15, 18 or 19 leaves.

At depth 4:

2 or 3 nodes and 3,4,5 or 8 leaves.

At depth 5 and 6:

Depth 5 is where most of the leaf variance is; there's a single (optional) node

```
[ 'Envelope' ] [ 'Payload-Metadata' ] [ 'HTTP-Response-Metadata' ] [ 'HTML-Metadata' ]
[ 'Head' ]
```

with 1,2,3,4,5 leaves under it at depth 6.

Let's eyeball a record with a larger number of json keys:

Took 0 sec. Last updated by anonymous at August 28 2017, 1:40:10 PM.

```
%pyspark
```

FINISHED

```
nkeys = 80
```

```
def f(rec):
    _,k = traverse(rec, 0, [])
    return len(k)

sample = files.\
    mapPartitionsWithSplit(json_mapper).\
    filter(lambda rec: f(rec) > nkeys).\
    take(100)

rec = sample[0]
depth, keys = traverse(rec, 0, [])
keys.sort()

print(len(keys))
for x in keys:
    print(x)
```

```
(5, 1, u'X-Passed-To-BeforeDispatch')
(5, 1, u'X-Passed-To-DLL')
(5, 1, u'X-Passed-To-PostProcessResponse')
(5, 1, u'X-Returned-From')
(5, 1, u'X-Returned-From-BeforeDispatch')
(5, 1, u'X-Returned-From-DLL')
(5, 1, u'X-Returned-From-PostProcessResponse')
(5, 1, u'X-Served-By')
(5, 1, u'X-UA-Device')
(5, 1, u'X-Varnish')
(5, 1, u'X-Varnish-beresp-grace')
(5, 1, u'X-Varnish-beresp-status')
(5, 1, u'X-Varnish-beresp-ttl')
(5, 1, u'x-stale')
(6, 1, u'Link')
(6, 1, u'Metas')
(6, 1, u'Scripts')
(6, 1, u'Title')
```

Took 54 sec. Last updated by anonymous at August 28 2017, 1:41:22 PM.

Let's look, at depth 5, at the Head node and a leaf Links, which contain relevant information about outgoing links:

Took 0 sec. Last updated by anonymous at August 28 2017, 1:44:00 PM.

```
%pyspark
from pprint import pprint

node = rec['Envelope']['Payload-Metadata']['HTTP-Response-Metadata']['HTML-Metadata']['Head']
_,k = traverse(node, 0, [])

for x in k:
    print(x)
    print(type(node[x[2]]))
    if type(node[x[2]]) == list: print(len(node[x[2]]))

pprint(node)
```

ork for Denver City Electric Motor Company as an electric motor repairman. ',
 u'name': u'twitter:description'},
 {u'content': u'Gregory Young', u'property': u'og:title'},
 {u'content': u'http://www.ardmoreite.com/article/20121123/OBITUARIES/12112990

```

4',
    u'property': u'og:url'},
    {u'content': u'article', u'property': u'og:type'},
    {u'content': u'Gregory Young, age 52, passed away on November 20, 2012, in Ardmore, Okla. Services will be held at 10 a.m., Saturday, November 24, 2012, in The Chapel at ',
    u'property': u'og:description'}]],
u'Scripts': [{u'path': u'SCRIPT@/src',
    u'url': u '//cdnjs.cloudflare.com/ajax/libs/modernizr/2.6.2/modernizr.min.js?20160411-3'},
    {u'path': u'SCRIPT@/src',
    u'url': u '//cdnjs.cloudflare.com/ajax/libs/jquery/2.1.4/jquery.min.js?20160411-3'},
    {u'path': u'SCRIPT@/src'}.

```

Took 0 sec. Last updated by anonymous at August 28 2017, 1:42:23 PM.

Links is a json leaf, but behaves like a node in that its value is a list of dicts:

FINISHED

Took 0 sec. Last updated by anonymous at August 28 2017, 1:45:03 PM.

```
%pyspark
```

READY

```

node = rec['Envelope']['Payload-Metadata']['HTTP-Response-Metadata']['HTML-Metadata']['Links',k = traverse(node, 0, [])

for x in node2:
    pprint.pprint(x)
{u'path': u'A@/href',
 u'text': u'Feedback',
 u'url': u'mailto:yournews@ardmoreite.com'}
{u'path': u'A@/href',
 u'text': u'Terms of Use',
 u'url': u'http://www.gatehousemedia.com/terms_of_use/'}
{u'path': u'A@/href',
 u'text': u'Privacy Policy',
 u'url': u'http://www.gatehousemedia.com/privacy-policy/'}
{u'path': u'A@/href',
 u'text': u'GateHouse Media Publications',
 u'url': u'http://www.gatehousemedia.com/our-markets/'}
{u'path': u'A@/href',
 u'target': u'_blank',
 u'title': u'The Daily Ardmoreite Facebook',
 u'url': u'http://www.facebook.com/pages/The-Ardmoreite/143707431292?ref=fb'}
{u'path': u'A@/href',
 u'target': u'_blank',

```

Finally, compare the distributions (over json records) of the numbers of 'url' values seen in a record for the two nodes just discussed:

Took 0 sec. Last updated by anonymous at August 28 2017, 1:47:45 PM.

```
%pyspark
```

FINISHED

```
import matplotlib.pyplot as plt
```



```

def compare_links(record):
    try:
        set1 = set([x['url'] for x in record['Envelope']['Payload-Metadata']['HTTP-Response-Links'])
        set2 = set([x['url'] for x in record['Envelope']['Payload-Metadata']['HTTP-Response-Links'])
        return [len(set1), len(set2)]
    except KeyError:
        return [0,0]


def get_link_counts(id_, iterator):
    for uri in iterator:
        file = unpack(uri)
        for record in file:
            if record['Content-Type'] == 'application/json':
                record = json.loads(record.payload.read())
                yield compare_links(record)

link_counts = files.mapPartitionsWithIndex(get_link_counts)
link_count = link_counts.collect()

llinks = [x[0] for x in link_count]
rlinks = [x[1] for x in link_count]

plt.hist(rlinks, bins=30, alpha=0.4, color='red')
plt.hist(llinks, bins=30, alpha=0.4, color='green')
plt.xscale('log')
plt.yscale('log')
plt.show()

```



Took 14 sec. Last updated by anonymous at August 28 2017, 1:56:18 PM. (outdated)

%pyspark

READY