

Bill 6 - building do...

FINISHED

Building domain features from WAT

What this notebook does:

Extracts domain string signatures and uses these to construct feature vectors for domains.

Took 0 sec. Last updated by anonymous at October 20 2017, 3:31:55 PM.

FINISHED

```
%pyspark

from __future__ import print_function

nfiles = 1024
inputURI = "s3://billsdata.net/CommonCrawl/domain_paths_from_%d_WAT_files/" % nfiles

domains_rdd = sc.textFile(inputURI).map(eval)

domain_uri_count = domains_rdd\
    .map(lambda x: [len(x['path_set']),
                    sum( [ len(uri[0]) for uri in x['path_set'] ] ),
                    sum([len(uri[0].encode('utf-8')) for uri in x['path_set'] ] )])\
    .aggregate((0, 0, 0, 0),
              lambda acc, value: (acc[0] + 1, acc[1] + value[0], acc[2] + value[1], acc[3] + value[2]),
              lambda acc1, acc2: (acc1[0] + acc2[0], acc1[1] + acc2[1], acc1[2] + acc2[2], acc1[3] + acc2[3]))

print("Nr domains: %15d" % domain_uri_count[0])
print("Nr page URIs: %13d" % domain_uri_count[1])
print("Nr URI chars: %13d" % domain_uri_count[2])
print("Nr URI bytes: %13d" % domain_uri_count[3])

Nr domains:          4285451
Nr page URIs:       420643485
Nr URI chars:     20332763778
Nr URI bytes:     20381326882
```

Took 1 min 55 sec. Last updated by anonymous at October 20 2017, 4:26:00 PM.

FINISHED

```
nfiles = 1
Nr domains: 99807
Nr page URIs: 1691436
Nr URI chars: 63947884
Nr URI bytes: 64199858

nfiles = 128
Nr domains: 1197074
Nr page URIs: 70897255
Nr URI chars: 3300575852
Nr URI bytes: 3308899496
```

nfiles = 1024
Nr domains: 4285451
Nr page URIs: 420643485
Nr URI chars: 20332763778
Nr URI bytes: 20381326882

Write to S3 a single string for all domains:

Took 0 sec. Last updated by anonymous at October 20 2017, 4:18:22 PM.

%pyspark

FINISHED

```
def domain_string(domain, path_set):
    """
    Takes domain and concatenates with path URIs separated by newlines.
    """
    out = domain + '\n' + '\n'.join(sorted([x[0] for x in list(path_set)])) + '\n\n\n'
    return out

domain_string_rdd = domains_rdd\
    .map(lambda x: domain_string(x['domain'], x['path_set']))
domain_string_rdd.cache()

outputURI = "s3://billsdata.net/CommonCrawl/domain_string_from_%d_WAT_files" % nfiles
codec = "org.apache.hadoop.io.compress.GzipCodec"
domain_string_rdd.saveAsTextFile(outputURI, codec)
```

PythonRDD[386] at RDD at PythonRDD.scala:48

Took 0 sec. Last updated by anonymous at October 20 2017, 4:30:16 PM. (outdated)

Cluster	nr files	nr domains	nr page URIs	nr chars	time
16 x m4.large	1	168k	1.8M	63.7M	6 sec
16 x m3.xlarge	128	2.6M	71.8M	3.26B	1 min 4 sec
16 x m4.large	128	2.6M	71.8M	3.26B	48 sec

FINISHED

To concatenate into a single gzip file (may need to mount extra local disk space):

```
$ aws s3 sync
s3://billsdata.net/CommonCrawl/domain_string_from_1024_WAT_files/ ./tmp
$ gunzip -c ./tmp/part*.gz | cat | gzip -c > ./tmp/big_domain_string_1024.gz
$ rm ./tmp/part* ./tmp/_SUCCESS
$ aws s3 sync ./tmp s3://billsdata.net/CommonCrawl/
$ rm -r ./tmp
```

Took 0 sec. Last updated by anonymous at October 20 2017, 4:42:33 PM.

%pyspark

FINISHED

```
for x in domain_string_rdd.takeSample(False, 10):
    print(x)
```

trezvvi64.ru
/2011/04/
/2011/05/
/2011/10/
/2011/12/

```
/2012/01/  
/2012/03/  
/2012/04/  
/2012/05/  
/2012/06/  
/2012/09/  
/2012/12/  
/2013/03/  
/2013/04/  
/2013/05/  
/2013/06/  
/2013/07/
```

Took 2 min 31 sec. Last updated by anonymous at October 20 2017, 4:32:54 PM.

```
%pyspark
```

FINISHED

```
def nonlatin_detector(str):  
    """  
    Computes the excess nr bytes over nr characters in a string.  
    """  
    N = len(str)  
    return float(len(str.encode('utf-8')))/N  
  
nonlatin_dist = domain_string_rdd.map(nonlatin_detector).collect()
```


Took 4 sec. Last updated by anonymous at October 20 2017, 4:33:19 PM.

```
%pyspark
```

FINISHED

```
import matplotlib.pyplot as plt  
  
nonlatin = [x for x in nonlatin_dist if x > 1.0]  
  
print("Nr domains: %8d" % len(nonlatin_dist))  
print("Nr non-latin: %6d" % len(nonlatin))  
print("Min non-latin score: %.10f" % min(nonlatin))  
  
plt.hist(nonlatin_dist, bins=50)  
plt.yscale("log")  
plt.show()
```

```
Nr domains: 4285451  
Nr non-latin: 51969  
Min non-latin score: 1.0000000237
```



Took 4 sec. Last updated by anonymous at October 20 2017, 4:33:31 PM.

```
%pyspark
threshold = 1.05

nonlatin_rdd = domain_string_rdd\
    .filter(lambda x: nonlatin_detector(x) > threshold)

print("Nr domains: %d" % nonlatin_rdd.count())
for x in nonlatin_rdd.takeSample(False, 10):
    print(x)
```

FINISHED

```
Nr domains: 12287
gateway-hotel.co.jp
/
www.voebb.de
/aDISWeb/app
/
korona-m.com
/images/cart.png
/images/facebook.png
/images/logo.png
/images/magnifier.png
/products/thumbs/bireni-fastatsi-100gr-2016-05-24-11-48-39.jpg
/products/thumbs/bireni-fastatsi-200gr-2016-05-24-11-49-35.jpg
/products/thumbs/parzhena-tsarevitsa-100-gr-2016-05-24-11-56-36.jpg
/products/thumbs/parzhena-tsarevitsa-200-gr-2016-05-24-11-57-14.jpg
/products/thumbs/pechen-sham-fastak-100-gr-2016-05-24-11-55-47.jpg
/products/thumbs/pechen-sham-fastak-200-gr-2016-05-24-11-56-11.jpg
/products/thumbs/pechen-sham-fastak-100gr-2016-05-24-11-50-16.jpg
```

Took 10 sec. Last updated by anonymous at October 20 2017, 4:44:21 PM.

```
%pyspark
outputURI = "s3://billsdata.net/CommonCrawl/domain_string_nonlatin_1pt05_12287_from_%d_WAT.
```

FINISHED

```
codec = "org.apache.hadoop.io.compress.GzipCodec"
nonlatin rdd.saveAsTextFile(outputURI, codec)
```

Took 6 sec. Last updated by anonymous at October 20 2017, 4:44:59 PM.

Now let's look at basic statistics of the path URI for a domain...

READY

%pyspark

FINISHED

```
import re
from math import log
from collections import Counter

def hx(i):
    """
    Normalised 2-char hex representation of 0-255
    """
    a = hex(i)[2:]
    if len(a)<2: a = ''.join(['0',a])
    return a
hexabet = [hx(x) for x in range(256)]

def depth(uri):
    return uri[:-1].count('/')

def length(uri):
    return len(uri) - uri.count('/')

def byte_count(str):
    out = dict([(x,0) for x in hexabet])
    ct = dict(Counter([c.encode('hex') for c in str.encode('utf-8')]))
    for k in out.keys():
        if k in ct.keys():
            out[k] += ct[k]
    out = [v[1] for v in sorted(out.iteritems(), key=lambda (k,v): k)]
    out = [float(x)/sum(out) for x in out]
    return out

def string_features_v1(str):
    """
    Coarse first version of a feature vector for a string.
    A placeholder for stronger versions.
    """
    N = float(length(str))
    if N==0: return None
    U = float(len(str.encode('utf-8')))/N
    D = depth(str)/N
    b = len(re.findall(r'\.', str))/N
    c = len(re.findall(r'-', str))/N
    d = len(re.findall(r'_', str))/N
    cap = len(re.findall(r'[A-Z]', str))/N
    num = len(re.findall(r'[0-9]', str))/N
    return [log(N), log(U), D, b, c, d, num, cap]

def string_features_v2(str):
    """
    Version 2: combine the byte distribution with the previous string statistics.
    """
    return byte_count(str) + string_features_v1(str)
```

Took 0 sec. Last updated by anonymous at October 20 2017, 5:04:01 PM.

FINISHED

[illegible]

Dots are URIs, colours are domains.

ERROR

```
Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-7418246470775012421.py", line 367, in <module>
    raise Exception(traceback.format_exc())
Exception: Traceback (most recent call last):
  File "/tmp/zeppelin_pyspark-7418246470775012421.py", line 360, in <module>
    exec(code, _zcUserQueryNameSpace)
  File "<stdin>", line 5, in <module>
  File "/usr/lib/spark/python/pyspark/rdd.py", line 479, in takeSample
    initialCount = self.count()
  File "/usr/lib/spark/python/pyspark/rdd.py", line 1041, in count
    return self.mapPartitions(lambda i: [sum(1 for _ in i)]).sum()
  File "/usr/lib/spark/python/pyspark/rdd.py", line 1032, in sum
```

```

return self.mapPartitions(lambda x: [sum(x)]).fold(0, operator.add)
File "/usr/lib/spark/python/pyspark/rdd.py", line 906, in fold
  vals = self.mapPartitions(func).collect()
File "/usr/lib/spark/python/pyspark/rdd.py", line 809, in collect
  port = self.ctx._jvm.PythonRDD.collectAndServe(self._jrdd.rdd())

```

Took 2 sec. Last updated by anonymous at October 20 2017, 5:05:52 PM.

```

[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 1222 samples in 0.001s...
[t-SNE] Computed neighbors for 1222 samples in 0.038s...
[t-SNE] Computed conditional probabilities for sample 1000 / 1222
[t-SNE] Computed conditional probabilities for sample 1222 / 1222
[t-SNE] Mean sigma: 0.000000
[t-SNE] KL divergence after 250 iterations with early exaggeration: 48.648224
[t-SNE] Error after 1000 iterations: 0.210676
[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 1222 samples in 0.008s...
[t-SNE] Computed neighbors for 1222 samples in 0.302s...
[t-SNE] Computed conditional probabilities for sample 1000 / 1222
[t-SNE] Computed conditional probabilities for sample 1222 / 1222
[t-SNE] Mean sigma: 0.080659
[t-SNE] KL divergence after 250 iterations with early exaggeration: 52.923714
[t-SNE] Error after 1000 iterations: 0.406656

```

READY

```

%pyspark

import matplotlib.pyplot as plt
for proj in [proj_2d_v1, proj_2d_v2]:
    fig, ax = plt.subplots(figsize=(12,6))
    cax = ax.scatter(proj[:,0], proj[:,1], s=10.0, c=col, edgecolors='face', cmap='rainbow')
    cbar = fig.colorbar(cax, ticks=range(ndomains), shrink=0.7)
    cbar.ax.set_yticklabels([dom[0] for dom in some_domains]) # vertically oriented color
    plt.show()


```

READY


```

[<matplotlib.text.Text object at 0x7f1d16015110>, <matplotlib.text.Text object at 0x7f1d160208d0>, <matplotlib.text.Text object at 0x7f1d15fdaad0>, <matplotlib.text.Text object at 0x7f1d15fd1610>, <matplotlib.text.Text object at 0x7f1d15fda450>, <matplotlib.text.Text object at 0x7f1d15fe4390>]

```



[, , , ,]



```
%pyspark
```

READY

```
page_feature_rdd.unpersist()
domains_rdd.unpersist()
```

PythonRDD[70] at RDD at PythonRDD.scala:48

READY

Export domain feature vectors

```
%pyspark
```

READY

```
nfiles = 1024
inputURI = "s3://billsdata.net/CommonCrawl/domain_paths_from_%d_WAT_files/" % nfiles
domains_rdd = sc.textFile(inputURI).map(eval)
domains_rdd.cache()

def domain_features(domain, path_set):
    """
    Takes domain + set of paths as output by parse_urls() and
    applies extracts statistics of the signature string.
    """
    return string_features_v2(domain_string(domain, path_set))

def feature_extractor(x):
    return (x['domain'], domain_features(x['domain'], x['path_set']))

domain_feature_rdd = domains_rdd.map(feature_extractor)
```

```
%pyspark
```

READY

```
outputURI = "s3://billsdata.net/CommonCrawl/domain_basic_string_feature_vectors_from_%d_WA"
codec = "org.apache.hadoop.io.compress.GzipCodec"
```



```
domain_feature_rdd.saveAsTextFile(outputURI, codec)
```

Timings:

READY

Cluster	nr files	nr domains	time
16 x m4.large	128	2.6M	40 min 7 sec

Let's check what we've just written:

```
%pyspark
```

READY

```
inputURI = "s3://billsdata.net/CommonCrawl/domain_hex_feature_vectors_from_%d_WAT_files" %
features_rdd = sc.textFile(inputURI).map(eval)
print("Nr domains:", features_rdd.count())
print(features_rdd.take(1))
```

[illegible]

```
%pyspark
```

READY