

Bill 6 - simple dom...

FINISHED

Building domain features from WAT

What this notebook does:

Extracts domain string signatures and uses these to construct feature vectors for domains.

Took 0 sec. Last updated by anonymous at October 15 2017, 11:59:35 AM.

```
%pyspark
```

FINISHED

```
domains_rdd.unpersist()
domain_string_rdd.unpersist()
```

PythonRDD[211] at RDD at PythonRDD.scala:48

Took 0 sec. Last updated by anonymous at October 15 2017, 3:18:10 PM.

```
%pyspark
```

FINISHED

```
from __future__ import print_function
```

```
nfiles = 1
```

```
inputURI = "s3://billsdata.net/CommonCrawl/domain_paths_from_%d_WAT_files/" % nfiles
```

```
domains_rdd = sc.textFile(inputURI).map(eval)
```

```
domain_uri_count = domains_rdd\
    .map(lambda x: [len(x['path_set']),
                    sum([len(uri) for uri in x['path_set']]),
                    sum([len(uri.encode('utf-8')) for uri in x['path_set'])])
    .aggregate((0, 0, 0, 0),
               lambda acc, value: (acc[0] + 1, acc[1] + value[0], acc[2] +
                                     value[1], acc[3] + value[2]),
               lambda acc1, acc2: (acc1[0] + acc2[0], acc1[1] + acc2[1],
                                     acc1[2] + acc2[2], acc1[3] + acc2[3]))
```

```
print("Nr domains: %15d" % domain_uri_count[0])
print("Nr page URIs: %13d" % domain_uri_count[1])
print("Nr URI chars: %13d" % domain_uri_count[2])
print("Nr URI bytes: %13d" % domain_uri_count[3])
```

```
Nr domains:          168033
Nr page URIs:        1782572
Nr URI chars:        63676121
Nr URI bytes:        63928095
```

Took 4 sec. Last updated by anonymous at October 15 2017, 3:18:24 PM.

```
nfiles = 1
```

FINISHED

```
Nr domains: 168033
```

```
Nr page URIs: 1782572
```

```
Nr URI chars: 63676121
```

```
Nr URI bytes: 63928095
```

nfiles = 128
Nr domains: 2626203
Nr page URIs: 71799497
Nr URI chars: 3259974688
Nr URI bytes: 3268298469

nfiles = 1024
Nr domains: 10802408
Nr page URIs: 420975127
Nr URI chars: 19980667843
Nr URI bytes: 20029236547

Write to S3 a single string for all domains:

Took 0 sec. Last updated by anonymous at October 14 2017, 8:10:18 PM.

%pyspark

FINISHED

```
def domain_string(domain, path_set):  
    """  
    Takes domain and concatenates with sorted path URIs separated by newlines.  
    """  
    out = domain + '\n' + '\n'.join(sorted(list(path_set))) + '\n\n\n'  
    return out  
  
domain_string_rdd = domains_rdd\  
    .map(lambda x: domain_string(x['domain'], x['path_set']))  
domain_string_rdd.cache()  
  
outputURI = "s3://billsdata.net/CommonCrawl/domain_string_from_%d_WAT_files" % nfiles  
codec = "org.apache.hadoop.io.compress.GzipCodec"  
domain_string_rdd.saveAsTextFile(outputURI, codec)
```

Took 0 sec. Last updated by anonymous at October 15 2017, 3:18:48 PM. (outdated)

Cluster	nr files	nr domains	nr page URIs	nr chars	time
16 x m4.large	1	168k	1.8M	63.7M	6 sec
16 x m3.xlarge	128	2.6M	71.8M	3.26B	1 min 4 sec
16 x m4.large	128	2.6M	71.8M	3.26B	48 sec

FINISHED

To concatenate into a single gzip file (may need to mount extra local disk space):

```
$ aws s3 sync s3://billsdata.net/CommonCrawl/domain_string_from_128_WAT_files/  
./tmp  
$ gunzip -c ./tmp/part*.gz | cat | gzip -c > ./tmp/big_domain_string_128.gz  
$ rm ./tmp/part* ./tmp/_SUCCESS  
$ aws s3 sync ./tmp s3://billsdata.net/CommonCrawl/  
$ rm -r ./tmp
```

Took 0 sec. Last updated by anonymous at October 14 2017, 12:10:53 PM.

%pyspark

FINISHED

```
for x in domain_string_rdd.takeSample(False, 10):  
    print(x)
```

softportal.com

```
/img/draugda.png
www.seksgratka.com
/
www.canadianprogressiveworld.com
search.sify.com
/
www.okcupid.com
/
/about
/careers
/legal/privacy
/legal/safety-tips
/legal/terms
/login
/mobile
/press
/profile/Escape520
```

Took 4 sec. Last updated by anonymous at October 15 2017, 3:19:00 PM.

```
%pyspark
```

FINISHED

```
def nonlatin_detector(str):
    """
    Computes the excess nr bytes over nr characters in a string.
    """
    N = len(str)
    return float(len(str.encode('utf-8')))/N

nonlatin_dist = domain_string_rdd.map(nonlatin_detector).collect()
```

Took 1 sec. Last updated by anonymous at October 15 2017, 3:19:10 PM.

```
%pyspark
```

FINISHED


```
import matplotlib.pyplot as plt

nonlatin = [x for x in nonlatin_dist if x > 1.0]

print("Nr domains: %8d" % len(nonlatin_dist))
print("Nr non-latin: %6d" % len(nonlatin))
print("Min non-latin score: %.10f" % min(nonlatin))

plt.hist(nonlatin_dist, bins=50)
plt.yscale("log")
plt.show()
```

```
Nr domains:   168033
Nr non-latin:    1154
Min non-latin score: 1.0000065218
```



Took 4 sec. Last updated by anonymous at October 15 2017, 3:19:17 PM.

```
%pyspark
threshold = 1.0

nonlatin_rdd = domain_string_rdd\
    .map(lambda s: [s, nonlatin_detector(s)])\
    .filter(lambda x: x[1] > threshold)

print("Nr domains: %d" % nonlatin_rdd.count())
for x in nonlatin_rdd.takeSample(False, 10):
    print(x[0])
```

FINISHED

```
Nr domains: 1154
ja.dbpedia.org
/resource/差別
satelonline.kz
/
/about_us
/adresa-servisnih-centrov
/akcia
/akkumulyatori-5
/akkumulyatori-5/aaaaa-palchikovie-10969
/akkumulyatori-5/acer-8422
/akkumulyatori-5/alcatel-8418
/akkumulyatori-5/alkalinovie-15538
/akkumulyatori-5/anydata-18763
/akkumulyatori-5/apple-8416
/akkumulyatori-5/asus-8423
/akkumulyatori-5/benqsiemens-7339
/akkumulyatori-5/blackberry-8417
```

Took 1 sec. Last updated by anonymous at October 15 2017, 3:19:33 PM.

```
%pyspark
```

FINISHED

```
outputURI = "s3://billsdata.net/CommonCrawl/domain_string_nonlatin_from_%d_WAT_files" % nf
codec = "org.apache.hadoop.io.compress.GzipCodec"
nonlatin rdd.saveAsTextFile(outputURI, codec)
```

Took 20 sec. Last updated by anonymous at October 15 2017, 11:10:17 AM.

Now let's look at basic statistics of the path URI for a domain...

READY

%pyspark

FINISHED

```
import re
from math import log
from collections import Counter

def hx(i):
    """
    Normalised 2-char hex representation of 0-255
    """
    a = hex(i)[2:]
    if len(a)<2: a = ''.join(['0',a])
    return a

hexabet = [hx(x) for x in range(256)]

def byte_count(str):
    out = dict([(x,0) for x in hexabet])
    ct = dict(Counter([c.encode('hex') for c in str.encode('utf-8')]))
    for k in out.keys():
        if k in ct.keys():
            out[k] += ct[k]
    out = [v[1] for v in sorted(out.iteritems(), key=lambda (k,v): k)]
    out = [float(x)/sum(out) for x in out]
    return out

def string_features_v1(str):
    """
    Coarse first version of a feature vector for a string.
    A placeholder for stronger versions.
    """
    N = float(len(str))
    if N==0: return None
    U = float(len(str.encode('utf-8')))

    a = len(re.findall(r'/', str))/N
    b = len(re.findall(r'\.', str))/N
    c = len(re.findall(r'-', str))/N
    d = len(re.findall(r'_', str))/N
    cap = len(re.findall(r'[A-Z]', str))/N
    num = len(re.findall(r'[0-9]', str))/N
    return [log(N), log(U), a, b, c, d, num, cap]

def string_features_v2(str):
    """
    Version 2: combine the byte distribution with the previous string statistics.
    """
    return byte_count(str) + string_features_v1(str)
```

Took 0 sec. Last updated by anonymous at October 15 2017, 3:20:24 PM.

%pyspark

FINISHED

```
def feature_extractor(x):
    str_set = [s for s in x['path_set'] if (string_features_v1(s) is not None) and (string.
a = [string_features_v1(s) for s in str_set]
b = [string_features_v2(s) for s in str_set]
    return (x['domain'], a, b)

page_feature_rdd = domains_rdd.map(feature_extractor)
page_feature_rdd.cache()
```

PythonRDD[249] at RDD at PythonRDD.scala:48

Took 0 sec. Last updated by anonymous at October 15 2017, 3:20:39 PM.

The plot below takes a random sample of domains, and computes feature vectors v1 and v2 from their path URIs for each domain.

Dots are URIs, colours are domains.

Took 0 sec. Last updated by anonymous at October 15 2017, 3:22:51 PM.

%pyspark

FINISHED

```
ndomains = 6
minpaths = 50

some_domains = page_feature_rdd\
    .filter(lambda x: len(x[1]) >= minpaths)\
    .takeSample(False, ndomains)
```

Took 4 sec. Last updated by anonymous at October 15 2017, 3:29:40 PM.

```
[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 1222 samples in 0.001s...
[t-SNE] Computed neighbors for 1222 samples in 0.038s...
[t-SNE] Computed conditional probabilities for sample 1000 / 1222
[t-SNE] Computed conditional probabilities for sample 1222 / 1222
[t-SNE] Mean sigma: 0.000000
[t-SNE] KL divergence after 250 iterations with early exaggeration: 48.648224
[t-SNE] Error after 1000 iterations: 0.210676
[t-SNE] Computing 91 nearest neighbors...
[t-SNE] Indexed 1222 samples in 0.008s...
[t-SNE] Computed neighbors for 1222 samples in 0.302s...
[t-SNE] Computed conditional probabilities for sample 1000 / 1222
[t-SNE] Computed conditional probabilities for sample 1222 / 1222
[t-SNE] Mean sigma: 0.080659
[t-SNE] KL divergence after 250 iterations with early exaggeration: 52.923714
[t-SNE] Error after 1000 iterations: 0.406656
```

FINISHED


Took 45 sec. Last updated by anonymous at October 15 2017, 3:30:27 PM.

%pyspark


FINISHED

```
import matplotlib.pyplot as plt
for proj in [proj_2d_v1, proj_2d_v2]:
    fig, ax = plt.subplots(figsize=(12,6))
    cax = ax.scatter(proj[:,0], proj[:,1], s=10.0, c=col, edgecolors='face', cmap='rainbow')
    cbar = fig.colorbar(cax, ticks=range(ndomains), shrink=0.7)
    cbar.ax.set_yticklabels([dom[0] for dom in some_domains]) # vertically oriented color
    plt.show()
```

[<matplotlib.text.Text object at 0x7f1d16015110>, <matplotlib.text.Text object at 0x7f1d160208d0>, <matplotlib.text.Text object at 0x7f1d15fdaad0>, <matplotlib.text.Text object at 0x7f1d15fd1610>, <matplotlib.text.Text object at 0x7f1d15fda450>, <matplotlib.text.Text object at 0x7f1d15fe4390>]



[, , , ,]



Took 4 sec. Last updated by anonymous at October 15 2017, 3:30:50 PM.

```
%pyspark
```

READY

```
page_feature_rdd.unpersist()  
domains_rdd.unpersist()
```

PythonRDD[70] at RDD at PythonRDD.scala:48

READY

Export domain feature vectors

READY

READY

READY

Let's check what we've just written:

READY

[illegible]

%pyspark

READY